


Twelve days into Beta, the foundation is built and merged.

Every open merge request on GitLab and GitHub landed on main today. The merged build passes all 1,365 automated tests. The numbers below come from the GitLab boards and the git history.

Commits per day on main

149 commits since Sept 21. Sept 22 was the biggest day, with 55.



74 of 223

Beta tickets closed or at Evidence, with 122 more ready to pull

39 of 40

merge requests merged on GitLab, 8 of them today

1,365

automated tests passing: 455 frontend, 910 Rust

21 + 23

written decision records and evidence files in the repo

Beta board, all four lanes



Anurup

46 tickets done. Search baseline, capture, memory pipeline

Kunj

14 tickets done. Command contract, tool registry, risk policy

Minh

3 reopen tickets done: QA matrix, unit tests, file paths

Felipe

10 tickets done. Voice system, Touch ID

The Alpha demo scored 96 out of 100.

Feedback from Shamit Fatin, received before this week's meeting.

96

out of 100 on the Alpha demo

Excellent work on a very strong, polished system.
Continue working on the application for next iteration.

Shamit Fatin, Alpha Demo grade and feedback notes

To do before this week's meeting

Every team member reads the grade and the feedback notes. The notes name no specific problems to fix, so we read "continue" as: keep deepening what the demo showed.

What we did with it

- Measured search before changing it: Recall@5 0.955 on the seeded week.
- Built one voice control, Touch ID without a runtime compiler, and a typed command registry.
- Put the next four weeks on a ticket board with evidence rules.

Project status: the foundation landed, Beta is about proof.

74

Beta tickets closed or at Evidence since Sept 21

Done in the past two weeks

Voice system built

Native speech helper, Rust session manager, permissions, and one shared control. Home uses it.

Touch ID works

Compiled helper, no Swift compiler needed at runtime.

Command surface

Contract, typed tool registry, and one risk policy for command bar, voice and MCP.

Measured baselines

Search and vault (VS-01), voice entry points (VO-01), one memory journey through the real pipeline (VS-27).

Reopen, notch and Agent

Reopen QA matrix and native file paths, notch HUD with spoken Do mode, Agent page as one chat.

39 of 40

merge requests merged, all open ones cleared today

Goals before end of Beta, Oct 21

Search Recall@5 at 0.90 or higher

Paraphrase queries at 0.80 or higher, one retrieval path for Search, Ask and MCP.

p95 search under 500 ms

Checked at 10,000 memories.

Exact reopen for 90% of a live day

Starts at 10% today.

Richer stored text

Median memory text from 129 to 800 characters, chunk vectors for every memory.

Voice commands

First partial text in 0.5 s, 45 of 50 scripted commands done right.

Five destinations

Home, Search and Ask, Vault, Daily Brief, Trust and Settings.

0.955

Search Recall@5 baseline, with Ask at 1.000

Roadblocks and help we want

One dependency between lanes

Search voice needs Kunj's command system before Felipe can finish it.

Small real dataset

The owner vault has 29 memories over 5 active days, so scale and long-run quality are unproven. Why most real memories carry little text is still being investigated (VS-14).

Outside users

We need people beyond the team for the user study.

help from staff?

8 GB hardware budget

Always-on CPU and memory targets must hold with the local model idle.

Reflective review: the process works, the board lagged the code.

What is working well

Tickets carry their own proof

212 tickets each say what to do, how to verify, and what evidence to attach.

Measure first, then change

23 evidence files and 21 decision records; search and voice both got baselines before any tuning.

Small stacked merge requests

Seven voice and onboarding MRs merged cleanly in order.

Automated safety net

1,365 tests pass on the merged build, and every GitHub pull request runs 8 automated checks.

Where we can improve

Board behind the code

Today we found finished work with no ticket moved: RE-02, RE-03, GS-11, and five pieces of shipped work with no ticket at all.

Review waits

Eight merge requests sat open together before being cleared today.

No CI on GitLab merge requests

None of the open MRs showed a pipeline, so a field added in one lane broke tests in another until we merged and ran them.

Work lands in bursts

55 commits on Sept 22, then quiet days.

Plans to improve

Move the ticket when you open the MR

Already step 6 of our workflow; the Wednesday written check-in (PD-03) makes it visible.

Reviewer closes within one working day

Response norm written into TEAM.md.

Merge gates in GitLab

VS-29 stages the quality checks so a merge request cannot land red.

Friday scoreboard

PD-05 prints retrieval, vault health, voice and reopen numbers every week.

User study: planned, not yet run.

No structured user-study session has been run yet.

Status

Six discovery conversations, all Ready

PD-18, PD-19 and PD-20: two 20-minute conversations each with knowledge workers outside the team, on getting back into work, quick actions, and refinding files. No FNDR demo until the end.

Five-user task study, planned

30 minutes each with a fixed Resume task, a System Usability Scale score, and quotes. Felipe leads with Anurup.

Success target

4 of 5 users finish the Resume task unaided.

Our thoughts

The one number we care about is time to correctly recover prior work, compared with the person's normal tools. The conversations come first so the task script reflects how people really resume work, not how we assume they do.

Questions for the course staff

draft, team to confirm

- Do sessions with classmates or outside volunteers need consent or review, since FNDR records on-screen content?
- Should participants use a seeded sample profile instead of their own Mac, to avoid collecting real data?
- Can the course help recruit five people outside the team?
- Is a SUS score plus task time enough, or do you want a different measure?
- Our plan runs the full study before the Final (Dec 14), with discovery conversations before Beta. Is that the right timing?

Anurup

Lead. Search and memory quality, capture, Screen Guide reliability, team operations

46

Beta tickets closed or at Evidence

104

commits, 34,452 code lines changed

68

issues closed all time

Beta tickets assigned: 83



■ Closed **44** ■ At Evidence **2** ■ Doing **3**

■ Ready **34**

FNDR Beta sprint progress report, Sept 21 to Oct 2, 2026

What shipped

Search has a measured baseline.

VS-01 recorded Recall@5 0.955 and MRR@10 0.909 for Search, and 1.000 for Ask, on a 20-record knowledge-worker week. Every later search change can now be judged against it.

One memory followed through the real pipeline.

VS-27 records a single memory from capture to result with event-backed evidence, and the six-journey baseline is staged behind it.

Capture became measurable and safer.

ScreenCaptureKit one-shot capture, perceptual-hash dedupe with loop detection, per-stage p50 and p95 timings, and a Privacy Proof screen showing what was skipped and why.

The memory pipeline was audited and covered.

MEM-01 to MEM-11: dead and duplicate code traced, finalization invariants under test, storage indexes, replay sessions with gold stories.

Screen Guide is steadier on Retina displays.

Capture handoff and Retina OCR fixes landed this week, plus an explicit private mode. Full-screen stabilization (GS-15) is in progress.

The team runs on a board.

GitLab labels, weekly milestones, issue and MR templates, TEAM.md, the demo script, and the agent-switch protocol all came from OPS-01 to OPS-12.

Next

Finish Screen Guide diagnostics (GS-15, GS-17), prove activity traces follow real process boundaries (PX-07), and run the six-journey human and agent baseline (VS-28).

Kunj

Command surface, agent, Screen Guide
delivery, release engineering

26

commits, 15,154 code lines changed

14

Beta tickets closed or at Evidence

6

merge requests merged

Beta tickets assigned: 45



■ Closed **3** ■ At Evidence **11** ■ Ready **31**

What shipped

The command bar has a contract.

GS-01 wrote the command-surface contract and amended ADR-014; PD-06 and ADR-022 define what the bar may do this semester.

One typed tool registry.

GS-03 gives every command a typed definition, so voice, the command bar and MCP all call the same tools.

One risk policy for every action.

GS-11 applies the same policy to the command bar, voice and MCP. It is merged to main and now at Evidence on the board.

FNDR can listen at the notch and act on the Mac (FEA-11).

The notch HUD takes spoken requests, and Do mode lets FNDR narrate and operate the Mac with an approval prompt on every action.

Screen Guide shipped natively.

Native delivery, a fixed-state status item, safe local voice input, and privacy-bounded filename lookup (FEA-07 to FEA-10).

The Agent page is one chat (FEA-12).

History, memory context, ChatGPT sign-in, and Hermes back as a model provider.

Releases open, and CI tells the truth (FEA-13, FEA-14).

The release pipeline now produces a DMG people can open. Two hidden Rust test failures, search latency budgets and a dependency bump were fixed so CI stops blocking the team.

Next

Connect Search voice to the command system so Felipe can finish the last Search voice work. This is the one dependency between the two lanes.

Minh

Reopen, embeddings, capture reliability

3

reopen tickets closed or at Evidence (RE-01, RE-02, RE-03)

11

merge requests merged all time

10

issues closed all time

Beta tickets assigned: 32



■ Closed **1** ■ At Evidence **2** ■ Ready **29**

What shipped

Reopen has a QA matrix and unit tests.

RE-01 built the matrix of reopen scenarios on today's build. RE-02 unit-tests how reopen targets are built and resolved, before any rule changes.

Reopen now keeps real file paths.

RE-03 stores native document paths so a memory can take you back to the actual file. It merged today after a conflict with the Screen Guide focus fix was resolved.

Capture runs on its own thread.

The capture pipeline moved to a dedicated thread with a larger stack, and IOKit pauses capture when the keyboard and mouse go idle.

Embedding work shipped earlier in the semester.

BGE v5 reindexing in bounded batches, a fix for BGE query chunk search, and the semantic ranking system for ordering results.

Next

Move the embedding model choice to measurement (VS-17), then use RE-03 paths to lift exact reopen coverage, which the vault baseline puts at 10.3% today.

Felipe

Voice, onboarding, product decisions

10

Beta tickets closed or at Evidence

7

merge requests landed today, in stack order

6

voice entry points inventoried, converging on one control

Beta tickets assigned: 38



■ Closed **4** ■ At Evidence **6** ■ Ready **28**

What shipped

The voice system is built.

Contract (VO-02), native speech helper (VO-03), and a Rust session manager (VO-04) that streams events to the UI.

Permissions and Touch ID work.

Speech and microphone permissions (VO-05). Touch ID now uses a compiled helper, so the app no longer needs the Swift compiler at runtime (OB-02).

One shared voice control for the whole app.

VO-06 built it to replace the separate voice paths. Home now uses it for normal voice input (VO-07), and Search is next.

Voice was measured before it was changed.

VO-01 inventoried every entry point and recorded smoke baselines for Home and Screen Guide.

The product decisions are written down.

PD-15 fixes the five destinations and Labs (ADR-023). PD-16 sets how voice behaves.

Search bar voice is in progress, per Felipe.

Transcription appears asynchronously and nothing searches until the arrow or Enter is used. This is not on main yet.

Next

Wire Search voice commands into Kunj's command system, then finish the last Search voice work. Blocked only on that connection.